

ABHISHEK MOHTA

(412) 320-9043 ◇ amohta@andrew.cmu.edu ◇ abhimohta.github.io ◇ github.com/abhimohta

EDUCATION

Carnegie Mellon University - School of Computer Science

December 2019

Masters in Computer Science

Coursework: Machine Learning*, Deep Learning*, Probability and Mathematical Statistics*, Computer Systems

Birla Institute of Technology and Science (BITS) Pilani, Goa, India

July 2017

M.Sc. Economics, B.E. Computer Science, CGPA: 9.41/10

Merit Scholar - 8/10 semesters (top 2% of >600 students)

SKILLS

Programming Languages: C, Java, Python, C#

Tools and Frameworks: PyTorch, CNTK, Git, Powershell, Visual Studio

EXPERIENCE

Microsoft Research, Research Fellow, Bangalore, India

July 2017 - June 2018

- Guide: Dr. Prateek Jain, Senior Researcher, Microsoft Research India
- Enhanced the PROSE framework (in MS Excel) with machine learned models replacing hand-written heuristics.
- Implemented Neural Guided Deductive Search (NGDS) - a hybrid of symbolic logic techniques and statistical models - using LSTMs in CNTK to prune branches in the synthesis process improving performance by >75%.
- Paper published at International Conference on Learning Representations (ICLR), 2018 and patent filed. [Link]

Microsoft Research, Research Intern, Bangalore, India

Jan 2017 - June 2017

- Enabled Whole Program Analysis as a service - built a scalable and reliable end-to-end system to statically analyze code and find bugs leveraging Static Module Verifier (SMV).
- Successfully ran on the Windows drivers code base and found 350+ bugs - null pointers and use-after-free.

Amazon Development Centre, SDE Intern, Bangalore, India

July 2016 - Dec 2016

- Created new features in the Automated Content Correction and Validation engine based on n-gram models.
- Implemented a part of the infrastructure for a feedback system to reduce manual intervention for data validation.

Microsoft Development Center, Summer Intern, Hyderabad, India

May 2016 - July 2016

- Implemented a classifier to find new/missing restaurants in Bing leveraging query logs, achieved >90% accuracy.

PROJECTS/ASSIGNMENTS (GITHUB) (KAGGLE)

Development of Deep learning models on speech WSJ dataset

Carnegie Mellon — Fall 2018

- Implemented the LAS architecture - speech to text using attention networks and beam search
- Phoneme prediction by learning Language Model using CNNs, LSTMs, and CTC loss
- Speaker embedding verification using CNNs and transfer learning; Speech frame classification using MLPs
- Implemented the backpropagation algorithm of MLPs from scratch, including momentum and batch normalization

Development of Machine Learning models

Carnegie Mellon — Fall 2018

- Implemented decision trees, logistic regression, kNNs, K-means, neural networks and basic HMMs

Development of integral systems components from scratch

Carnegie Mellon — Summer 2018

- Implemented the malloc/calloc/free function, a proxy server with cache, a linux shell (tsh-tiny shell) and a generic cache simulator

Clustering terrorist attacks with optimized K-Means algorithm

BITS Pilani — Spring 2016

- Built clusters of terrorist attacks leveraging the Global Terrorism database using Hadoop and Java
- Clusters based on location, casualties, target group, etc. were created, which closely resembled the ground truth.

Independent Projects

2015 - 2017

- Developed web portal to computerize management of manufacturing firm - handles >10 factories daily.
- Co-Founded Xinger - a fully automated online food ordering portal with a client base of >7000 and revenues >\$25,000 in the 2 years of operation. Handled end-to-end development and operations.